


# CS294-158 Deep Unsupervised Learning

## Lecture 3b: Likelihood Models III: Latent variable models





Pieter Abbeel, Peter Chen, Jonathan Ho, Aravind Srinivas

UC Berkeley

# Latent Variable Models

- Autoregressive models + Flows
  - All random variables are observed
- Latent Variable Models (LVMs):
  - Some random variables are hidden - we do not get to observe



# Why Latent Variable Models?

- Simpler, lower-dimensional representations of data often possible
  - Latent variable models hold the promise of automatically identifying those hidden representations

Obj1 @ (x,y)  
= Corgi, red  
& white

Obj2 @ (x,y)  
= Corgi, red  
& white,  
floppy left  
ear



Background  
= Wood  
bench in a  
park

Obj3 @ (x,y)  
= Corgi,  
Tri-color

# Why Latent Variable Models?

- AR models are slow to sample because all pixels (observation dims) are assumed to be dependent on each other
- We can make part of observation space independent *conditioned on some latent variables*
  - Latent variable models *can* have faster sampling by exploiting statistical patterns


# Latent Variable Models

- Sometimes, it's possible to design a latent variable model with an understanding of the causal process that generates data
- In general, we don't know what are the latent variables and how they interact with observations
  - Most popular models make little assumption about what are the latent variables
  - Best way to specify latent variables is still an active area of research

# A simple latent variable model

$$z = (z_1, z_2, \dots, z_K) \sim p(z; \beta) = \prod_{k=1}^K \beta_k^{z_k} (1 - \beta_k)^{1-z_k}$$

$$x = (x_1, x_2, \dots, x_L) \sim p_\theta(x|z) \Leftrightarrow \text{Bernoulli}(x_i; \text{DNN}(z))$$



# Training LVM

- Now we have a model with parameters, how to train it?
- Maximum Likelihood again!
- Even for small  $K = 32$ , it's a summation over 4B terms

$$\log p(x) = \log \left( \sum_z p(x|z)p(z) \right)$$

$$z = (z_1, z_2, \dots, z_K)$$

$$p(z; \phi) = \prod_{k=1}^K \phi_k^{z_k} (1 - \phi_k)^{1-z_k}$$

$$\theta \leftarrow \operatorname{argmax}_{\theta} \left[ \log p_{\theta}(x) = \log \left( \sum_z p_{\theta}(x|z)p(z; \phi) \right) \right]$$

# Variational Inference (VI)

- Recap: we have  $O(1)$  access to  $p(z)$  and  $p(x|z)$  and want to efficiently compute  $p(x)$
- Assume we have an oracle:  $p(z|x)$ 
  - intuition: what was the  $z$  that generates this  $x$ ?

$$p(x) = \frac{p(x|z)p(z)}{p(z|x)}$$

- But how do we obtain  $p(z|x)$ ?




# Variational Inference (VI)

Given this simple graphical model, we are interested in the posterior distribution of the latent variables.


$$p(z|x) = \frac{p(x, z)}{p(x)} = \frac{p(x, z)}{\sum_z p(x|z)p(z)}$$

Intractable



# Variational Inference (VI)

Key idea in variational inference is to approximate the true posterior with a variational distribution. This results in an optimization problem that minimizes the distance between these two distributions, ex: KL divergence.



# Variational Inference (VI)

- Let  $q_x(z)$  be our approximation of true posterior  $p(z|x)$ 
  - Then we pick a divergence to minimize between these two distributions

$$\begin{aligned} D_{\text{KL}} [q_x(z) \parallel p(z|x)] &= \mathbb{E}_{z \sim q_x(z)} [\log q_x(z) - \log p(z|x)] \\ &= \mathbb{E}_{z \sim q_x(z)} \left[ \log q_x(z) - \log \frac{p(z, x)}{p(x)} \right] \\ &= \mathbb{E}_{z \sim q_x(z)} [\log q_x(z) - \log p(z) - \log p(x|z) + \log p(x)] \\ &= \underbrace{\mathbb{E}_{z \sim q_x(z)} [\log q_x(z) - \log p(z) - \log p(x|z)]}_{\text{Only this part depends on } z} + \log p(x) \end{aligned}$$

- note: the expectation can be approximated by stochastic samples, and every term in expectation can be computed in  $O(1)$  now

# Variational Lower Bound (VLB)

- We now have an objective amenable to stochastic optimization

$$D_{\text{KL}} [q_x(z) \parallel p(z|x)] = \mathbb{E}_{z \sim q_x(z)} [\log q_x(z) - \log p(z) - \log p(x|z)] + \log p(x)$$

- Turns out we can get more out of this exercise

$$\begin{aligned} \log p(x) &= -\mathbb{E}_{z \sim q_x(z)} [\log q_x(z) - \log p(z) - \log p(x|z)] + D_{\text{KL}} [q_x(z) \parallel p(z|x)] \\ &= \underbrace{\mathbb{E}_{z \sim q_x(z)} [\log p(z) + \log p(x|z) - \log q_x(z)]}_{\text{Variational Lower Bound}} + \underbrace{D_{\text{KL}} [q_x(z) \parallel p(z|x)]}_{\geq 0} \end{aligned}$$

- note: the optimal  $q_x(z)$  of VLB is  $p(z|x)$ , at which point VLB is tight ( $= \log p(x)$ )

# VLB Maximization

- Given a data distribution  $\mathbf{x} \sim p_{\text{data}}$ , we can know train the generative model by maximizing the VLB under data distribution

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \left[ \mathbb{E}_{\mathbf{z} \sim q_x(\mathbf{z})} [\log p(\mathbf{z}) + \log p(\mathbf{x}|\mathbf{z}) - \log q_x(\mathbf{z})] \right] \\ & \leq \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log p(\mathbf{x})] \end{aligned}$$

# Variational Inference

- Core idea: learn an approximation of the intractable posterior
- Widely applicable:
  - learning latent variable generative models
  - estimation / maximization of Mutual Information
  - learnable dequantization
  - ....

# Stochastic optimization of VLB

$$z = (z_1, z_2, \dots, z_K) \sim p(z; \beta) = \prod_{k=1}^K \beta_k^{z_k} (1 - \beta_k)^{1-z_k}$$

$$x = (x_1, x_2, \dots, x_L) \sim p_\theta(x|z) \Leftrightarrow \text{Bernoulli}(x_i; \text{DNN}(z))$$

$$\text{VLB} = \mathbb{E}_{z \sim q(z; \phi(x))} [\log p(x|z) - \log q(z; \phi(x)) + \log p(z)]$$

In the Bernoulli setting, this becomes:

$$\mathbb{E}_{z \sim q(z; \phi(x))} [\log p(x|z; \theta) - \log q(z; \phi(x)) + \log p(z; \beta)]$$

Core problem: how to optimize the expectation from which  $z$  is drawn

# Wake-Sleep algorithm

- Note: VLB was derived from  $\min \text{KL}[q_x(z) || p(z|x)]$ , hard to optimize because  $z$  drawn from  $q_x(z)$
- What if we instead minimize  $\text{KL}[p(z|x) || q_x(z)]$  for any given  $x$ 
  - still hard because we don't know  $p(z|x)$
- Trick: we know  $z$  if we generate them!  $\mathbf{z} \sim p_\beta(\mathbf{z}), \mathbf{x} \sim p_\theta(\mathbf{x}|\mathbf{z})$ 
  - caveat  $\mathbf{x} \sim p_{\text{model}}$  instead of  $\mathbf{x} \sim p_{\text{data}}$
- Problem?
  - VLB is maximized with  $x$  drawn from  $p_{\text{data}}$
  - minimizing  $\text{kl}(p || q)$  with  $x$  drawn from  $p_{\text{model}}$  doesn't guarantee the bound is tight



# Wake-Sleep algorithm

## Wake Phase

- Sample  $x \sim p_{\text{data}}, z \sim q(z; \phi(x))$ .
- Maximize VLB with respect to  $\theta, \beta$ .

## Sleep Phase (model *dreaming* samples)

- Sample  $z \sim p(z; \beta), x \sim p(x|z; \theta)$ .
- Minimize  $KL(p(z|x)||q(z; \phi(x)))$  with respect to  $\phi$  now that we have samples from  $p_{\text{model}}$ .

## Reverse KL


$$KL(p(z|x)||q(z; \phi(x))) = E_{z \sim p(z|x)} \left[ \underbrace{\log p(z|x)}_{\text{independent of } \phi} - \log q(z; \phi(x)) \right]$$

# Amortized Inference

- $q(z; \phi(x))$ 
  - Not scalable to large dataset
  - Expensive to evaluate new data points

- Amortized Inference

$$VLB = \mathbb{E}_{z \sim q_\phi(\cdot|x)} [\log p_\theta(x|z) - \log q_\phi(z|x) + \log p(z)]$$



# Helmholtz Machine

- Helmholtz Machine [Dayan, P., Hinton, G. E.,... 1995]
  - Bernoulli latent code + observation space
  - Learned with Wake-sleep algorithm
- Did not scale to solve more complex problems due to limitations of wake-sleep

# Directly optimizing VLB

- Wake-Sleep not effective, especially when  $p_{\text{model}}$  is far away from  $p_{\text{data}}$
- Can we directly optimize VLB?

Recall, we want

$$\phi, \theta \leftarrow \operatorname{argmax}_{\theta, \phi} \left( \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\log p_{\theta}(x|z) - \log q_{\phi}(z|x) + \log p(z)] \right)$$

Optimization with respect to  $\phi$  is of the form

$$\operatorname{argmax}_{\phi} \mathbb{E}_{z \sim q_{\phi}} [f(z)]$$

Well studied problem in reinforcement learning where no assumption on  $f$  is made.

# Likelihood Ratio Estimator

We are interested in  $\operatorname{argmax}_{\phi} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [f(z)]$

How do we compute  $\nabla_{\phi} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [f(z)]$  ?

$$\nabla_{\phi} \sum_z q_{\phi}(z|x) f(z) = \sum_z \nabla_{\phi} q_{\phi}(z|x) f(z) = \sum_z \underbrace{\frac{\nabla_{\phi} q_{\phi}(z|x)}{q_{\phi}(z|x)}} f(z) q_{\phi}(z|x)$$

$$\Rightarrow \nabla_{\phi} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [f(z)] = \sum_z (\nabla_{\phi} \log q_{\phi}(z|x) f(z)) q_{\phi}(z|x) = \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\nabla_{\phi} \log q_{\phi}(z|x) f(z)]$$

$$\phi \leftarrow \phi + \alpha \nabla_{\phi} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\nabla_{\phi} \log q_{\phi}(z|x) f(z)]$$

**Issue:** High variance gradients, needs many samples of  $z$  to form a good estimate [[Demo](#)]

# Pathwise Derivative (PD)

One other way to optimize this objective when  $z$  is continuous is to cast  $z$  as a function of a simple fixed noise such as standard gaussian.

$$z = g(\epsilon, \phi), \epsilon \sim \mathcal{N}(0, I)$$

$$\mathbb{E}_{z \sim q_\phi(\cdot|x)} [f(z)] = \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [f(g(\epsilon, \phi))]$$

When  $f$  is differentiable,

$$\nabla_\phi \mathbb{E}_{z \sim q_\phi(\cdot|x)} [f(z)] = \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\nabla_\phi f(g(\epsilon, \phi))]$$

# Pathwise Derivative (PD)

- Stochastic gradient possible if  $z$  is continuous now (more technical condition?)
  - Common choice:  $\epsilon \sim \text{Normal}$ ,  $f(\epsilon) = \mu + \sigma \epsilon$
  - Any flow that you just learned!
- Also known as reparameterization trick
- Can work with only 1~2 samples
- Demo

# PD applied to VI

## Variational AutoEncoder

$q_\phi(z|x)$  is modeled as a Gaussian with parameters  $\mu$  and  $\sigma$  a DNN encoder (parameters  $\phi$ ) of  $x$ . The DNN decoder  $p_\theta(x|z)$  is differentiable.

$$\text{Let } z = \Sigma^{1/2}(x; \phi)\epsilon + \mu(x; \phi)$$

$$\begin{aligned}\text{VLB} &= \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\log p_\theta(x|z) - \log q_\phi(z|x) + \log p(z)] \\ &= \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\log p_\theta(x|z)] - KL(q_\phi(z|x) || p(z))\end{aligned}$$

$\nabla_\theta [\text{VLB}]$  and  $\nabla_\phi [\text{VLB}]$  can now be efficiently computed with SGD.




# Why is it called an autoencoder?

- We have seen that a variational autoencoder is a latent variable model with Gaussian prior  $p(z)$  and approximate posterior  $q(z|x)$ .
  - Why is it called an “autoencoder”?


$$\log p_{\theta}(x) \geq \underbrace{\left( E_{z \sim q_x(z)} \log p_{\theta}(x|z) \right)}_{\text{Reconstruction loss}} - \underbrace{KL(q_{\phi}(z|x) || p(z))}_{\text{Regularization}}$$

$L(\theta, \phi)$  - VAE objective


# VAE



# VAE

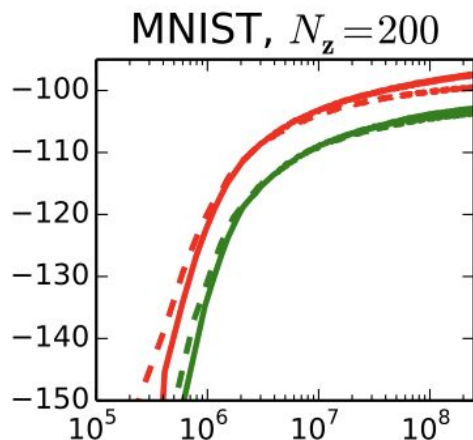


# VAE



# Compared to AR

- We now have a family of trainable latent variable models!
- But performance is lacking



Model	NLL Test
DBM 2hl [1]:	$\approx 84.62$
DBN 2hl [2]:	$\approx 84.55$
NADE [3]:	88.33
EoNADE 2hl (128 orderings) [3]:	85.10
EoNADE-5 2hl (128 orderings) [4]:	84.68
DLGM [5]:	$\approx 86.60$
DLGM 8 leapfrog steps [6]:	$\approx 85.51$
DARN 1hl [7]:	$\approx 84.13$
MADE 2hl (32 masks) [8]:	86.64
DRAW [9]:	$\leq 80.97$
PixelCNN:	81.30

# Bridging the gap between AR and VAE

- Improve Variational Inference
- Flexible decoder
- More expressive model architectures